

1. (2.5%) 訓練一個 model。

a. Generator: 與 colab 相同

Layer (type)	Output Shape	Param #
Linear-1	[-1, 8192]	819,200
BatchNorm1d-2	[-1, 8192]	16,384
ReLU-3	[-1, 8192]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	3,276,800
BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	819,200
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	204,800
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	4,803
Tanh-14	[-1, 3, 64, 64]	0
Total params: 5,142,083		
Trainable params: 5,142,083		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 19.62		
Estimated Total Size (MB): 22.62		

Discriminator: 與 colab 相同

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	4,864
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	204,928
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	819,456
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	3,277,312
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,193
Sigmoid-13	[-1, 1, 1, 1]	0
Total params: 4,316,545		
Trainable params: 4,316,545		
Non-trainable params: 0		
Input size (MB): 0.05		
Forward/backward pass size (MB): 2.31		
Params size (MB): 16.47		
Estimated Total Size (MB): 18.83		

Loss Function: BCELoss 與 colab 相同

Dataset: 助教提供

Optimizer: Adam, lr = 1e-4, betas = (0.5,0.999)

Epoch: 10

b. (1.5%) 請畫出至少 16 張 model 生成的圖片。



2. (3.5%) 請選擇下列其中一種 model : WGAN, WGAN-GP, LSGAN, SNGAN

(不要和 1. 使用的 model 一樣 , 至少 architecture 或是 loss function 要不同)

a. Generator: 將 ReLU 改成 LeakyReLU

Layer (type)	Output Shape	Param #
Linear-1	[-1, 8192]	819,200
BatchNorm1d-2	[-1, 8192]	16,384
LeakyReLU-3	[-1, 8192]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	3,276,800
BatchNorm2d-5	[-1, 256, 8, 8]	512
LeakyReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	819,200
BatchNorm2d-8	[-1, 128, 16, 16]	256
LeakyReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	204,800
BatchNorm2d-11	[-1, 64, 32, 32]	128
LeakyReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	4,803
Tanh-14	[-1, 3, 64, 64]	0
Total params: 5,142,083		
Trainable params: 5,142,083		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 19.62		
Estimated Total Size (MB): 22.62		

Discriminator: 與 colab 相同

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	4,864
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	204,928
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	819,456
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	3,277,312
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,193
Sigmoid-13	[-1, 1, 1, 1]	0

=====
Total params: 4,316,545
Trainable params: 4,316,545
Non-trainable params: 0
=====

Input size (MB): 0.05
Forward/backward pass size (MB): 2.31
Params size (MB): 16.47
Estimated Total Size (MB): 18.83
=====

Loss Function: BCELoss + 0.1 倍的 compute_gradient_penalty

```
def compute_gradient_penalty(D, real_samples, fake_samples):
    """Calculates the gradient penalty loss for WGAN GP"""
    # Random weight term for interpolation between real and fake samples
    alpha = Tensor(np.random.random((real_samples.size(0), 1, 1, 1)))
    # Get random interpolation between real and fake samples
    interpolates = (alpha * real_samples + ((1 - alpha) * fake_samples)).requires_grad_(True)
    d_interpolates = D(interpolates)


    fake = Variable(Tensor(real_samples.shape[0]).fill_(1.0), requires_grad=False)
    # Get gradient w.r.t. interpolates
    gradients = torch.autograd.grad(
        outputs=d_interpolates,
        inputs=interpolates,
        grad_outputs=fake,
        create_graph=True,
        retain_graph=True,
        only_inputs=True,
    )[0]
    gradients = gradients.view(gradients.size(0), -1)
    gradient_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean()
    return gradient_penalty
```

Dataset: 助教提供

Optimizer: Adam, lr = 1e-4, betas = (0.5, 0.999)

Epoch: 20

b. (1.5%) 和 1.b 一樣，就你選擇的 model，畫出至少 16 張 model 生成的圖片。



c. (1%) 請簡單探討你在 1. 使用的 model 和 2. 使用的 model，他們分別有何性質，描述你觀察到的異同。

第一題的 model 所畫出的圖片比較模糊，人的輪廓也有點奇怪，比方說眼睛的形狀和位置都會有點畸形。第二題的圖片就比較清晰，人的輪廓也較第一題正常且清晰許多。推測可能為計算 loss 的差異 (gradient penalty)

3. (4%) 請訓練一個會導致 mode collapse 的 model。

a. Generator: 與 colab 相同

Layer (type)	Output Shape	Param #
Linear-1	[-1, 8192]	819,200
BatchNorm1d-2	[-1, 8192]	16,384
ReLU-3	[-1, 8192]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	3,276,800
BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	819,200
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	204,800
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	4,803
Tanh-14	[-1, 3, 64, 64]	0
Total params: 5,142,083		
Trainable params: 5,142,083		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 19.62		
Estimated Total Size (MB): 22.62		

Discriminator: 與 colab 相同

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	4,864
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	204,928
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	819,456
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	3,277,312
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,193
Sigmoid-13	[-1, 1, 1, 1]	0
Total params: 4,316,545		
Trainable params: 4,316,545		
Non-trainable params: 0		
Input size (MB): 0.05		
Forward/backward pass size (MB): 2.31		
Params size (MB): 16.47		
Estimated Total Size (MB): 18.83		


Loss Function: BCELoss 與 colab 相同

Dataset: 助教提供

Optimizer: Adam, lr = 1e-4, betas = (0.5,0.999)


Epoch: 29

- b. (1.5%) 請畫出至少 16 張 model 生成且具有 mode collapse 現象的圖片。



- c. (1.5%) 在不改變 optimizer 和訓練 step 數的情況下，請嘗試使用一些方法來減緩 mode collapse。說明你嘗試了哪些方法，請至少舉出一種成功改善的方法，若有其它失敗的方法也可以記錄下來。

使用 WGAN_GP 的方法，在基礎條件不更動的情況下，生成的圖片 依然具有多樣性。效果如下：



除了將 Loss Function 改成 BCELoss + 0.1 倍的
compute_gradient_penalty , 其餘都和 3-a 相同